

---

# APACHE SPARK EXECUTION MODEL

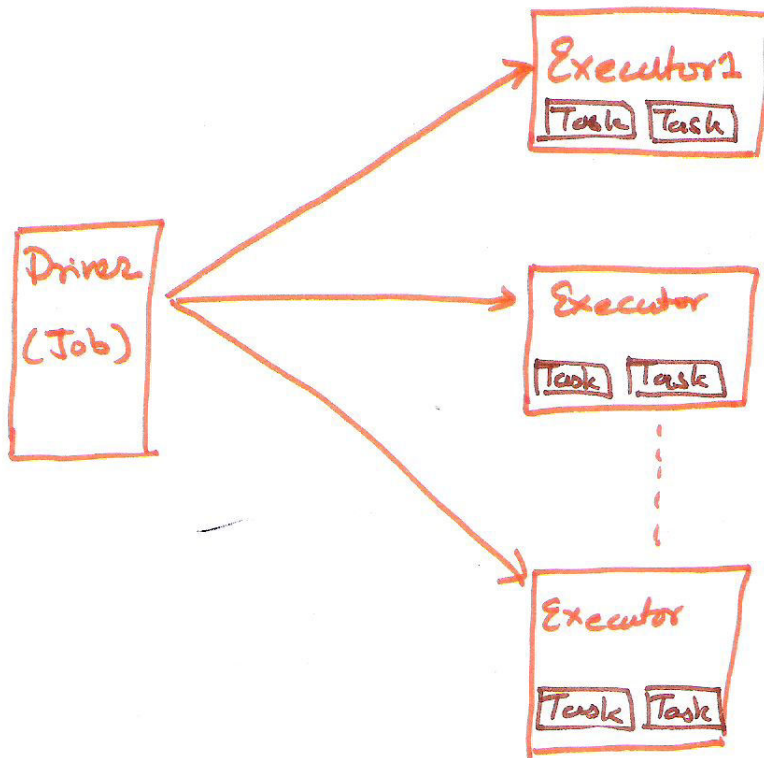
---

By [www.HadoopExam.com](http://www.HadoopExam.com)

**Note: These instructions should be used with the HadoopExam Apache Spark: Professional Trainings.  
Where it is executed and you can do hands on with trainer.**

## Spark Execution Model

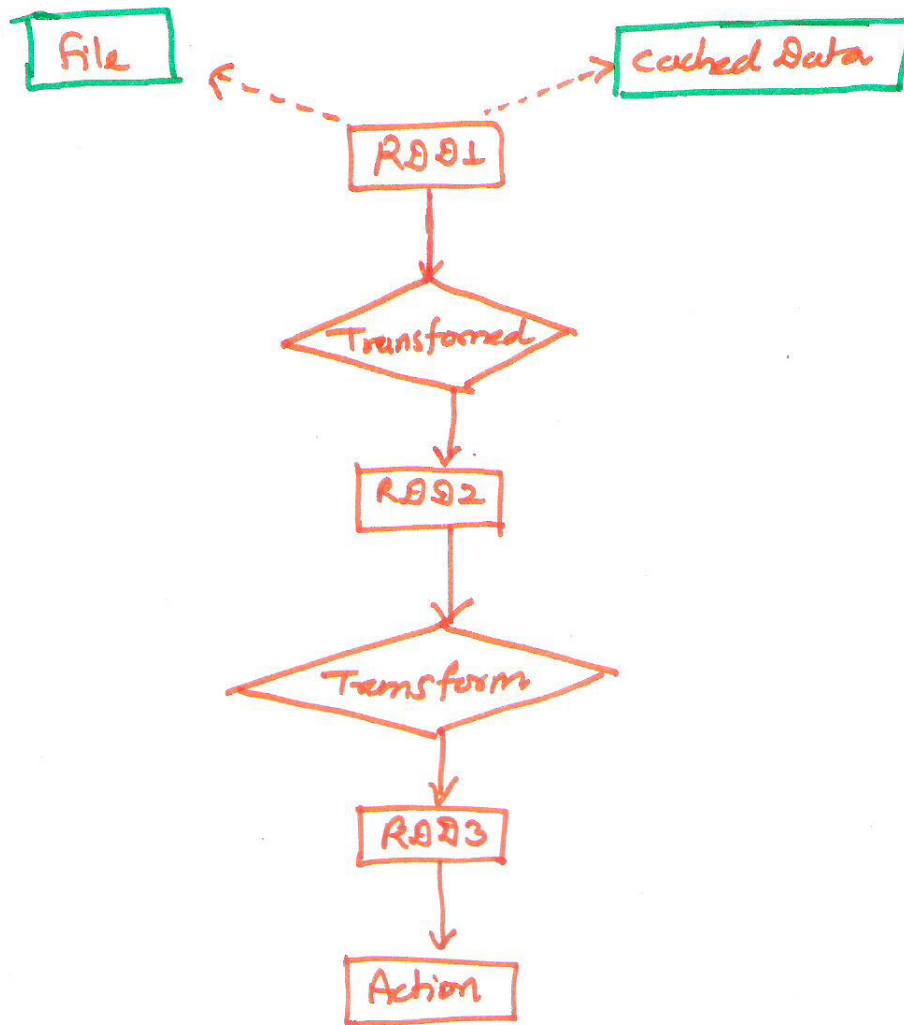
- How spark execute your programs.
- A spark Application consist of
  - Single Driver Process
  - Set of executor process [**Scattered across nodes**]
- Driver:- Control high level flow of work that needs to be done.
- Executor:- Executing the control flow, in the form of tasks
  - As well as storing the data, that the user choose to cache.




A Job contains tasks, and execute Concurrently.

- A single executor has a number of slots for running tasks, and will run many Concurrently.
- At the top of the execution hierarchy are jobs.
- Spark looks the graph of RDDs on which that action depends and formulate an execution plan.

- This plan starts with the farthest-back RDDs - that is, those that depends on no other RDDs.
- OR reference already cached-data and culminate in the final RDD required to produce the actions results.



- The execution plan consists of assembling the job's transformations into stages.






- A stage corresponds to a collection of tasks that all ~~execution~~ execute the same code, each on a different subset of data.

- Each stage contains a sequence of transformations that can be completed without shuffling the full data.

⇒ How & what determines whether data needs to be shuffled?

- RDD comprises a fixed number of partitions.



- Narrow transformation

- map
  - filter

- Work on partition, independently, no need data shuffling

- Spark also supports transformations with wide dependencies
- such as

- groupByKey()

- reduceByKey()


(4)

- This operation bring all the same key records from all the partition together to work upon.
- Shuffle transfer data around the cluster and results in a new stage with a new set of partitions.

```
sc.textFile("hadoopExam.log")
```

- map(mapfunc)
- flatMap(...)
- filter(...)
- count()

(Shuffle not Required)



- It executes a single action, which depends on a sequence of transformations on an RDD derived from a text file.
- It would executes stage concurrently on each all partitions.

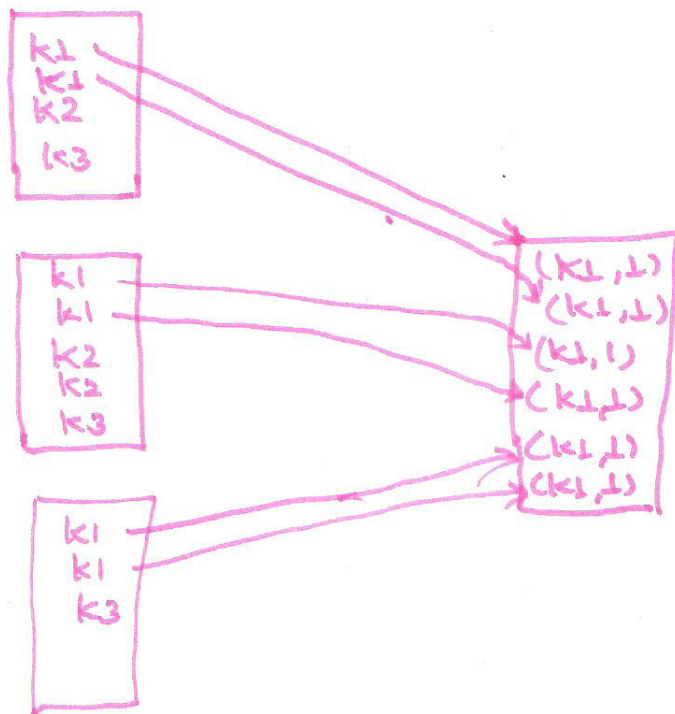


⇒ Because none of the output of these three operations depend on data that can come from different partitions than their input.

⇒ When shuffling happen?

```
Val token = sc.textFile("hdfs://hadoopexam.log")  
token.flatMap(_.split(' '))  
  • map(—, 1)  
  • reduceByKey(— + —) // Word Count  
  • flatMap(—, 1).toCharArray  
  • map(—, 1)  
  • reduceByKey(— + —) // char Count  
  • collect()
```

⇒ This process will run in more than one stages.

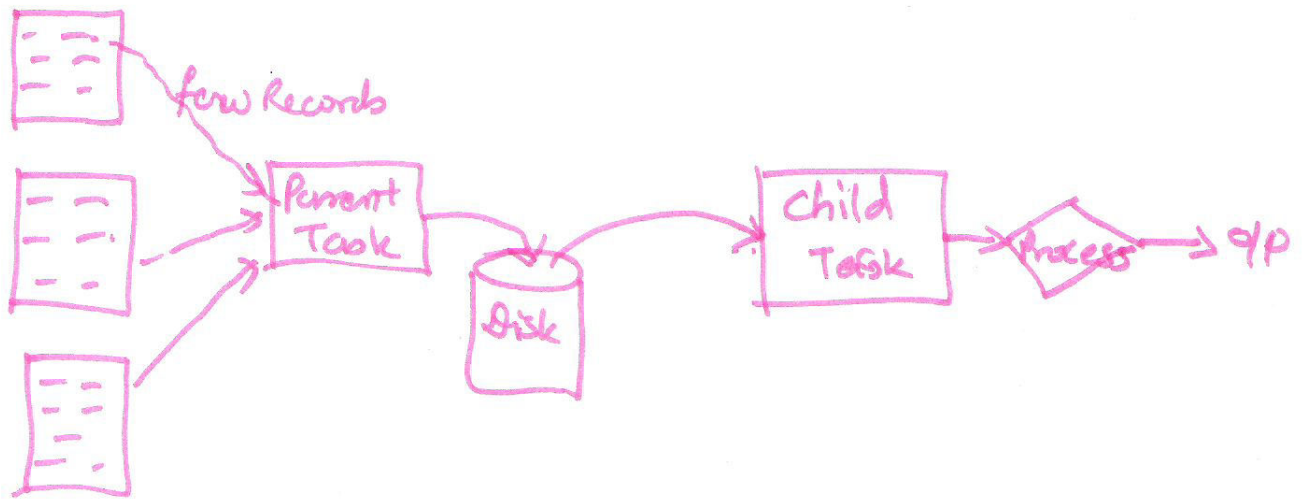


reduceByKey()  
require to be  
shuffled across  
partitions for getting  
same key

(k1, 6) → stage 1 (Word Count)  
→ stage 2 (character count)

⇒ At each stage boundary data is written to disk by task in parent stage.

⇒ Then fetched over the n/w by tasks in the child stage.



⇒ Because shuffle incur heavy disk and n/w I/O, stage boundaries can be expensive and should be avoided when possible.

⇒ The number of data partitions in parent stage may be different than the number of partitions in the child stage.

⇒ Transformations that may trigger a stage boundary typically accept "numPartitions" argument that determine how many partitions to split this the data into child stage.